

```
In [1]: 1 import pandas as pd
```

```
In [2]: 1 import numpy as np
```

```
In [3]: 1 df=pd.read_csv("C:\\Users\\cozze\\OneDrive\\Desktop\\nyc_rolling_sales.csv")
```

```
In [4]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 84548 entries, 0 to 84547
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            84548 non-null  int64
1   BOROUGH                               84548 non-null  int64
2   NEIGHBORHOOD                          84548 non-null  object
3   BUILDING CLASS CATEGORY               84548 non-null  object
4   TAX CLASS AT PRESENT                  84548 non-null  object
5   BLOCK                                 84548 non-null  int64
6   LOT                                   84548 non-null  int64
7   EASE-MENT                             84548 non-null  object
8   BUILDING CLASS AT PRESENT             84548 non-null  object
9   ADDRESS                               84548 non-null  object
10  APARTMENT NUMBER                      84548 non-null  object
11  ZIP CODE                              84548 non-null  int64
12  RESIDENTIAL UNITS                     84548 non-null  int64
13  COMMERCIAL UNITS                      84548 non-null  int64
14  TOTAL UNITS                           84548 non-null  int64
15  LAND SQUARE FEET                     84548 non-null  object
16  GROSS SQUARE FEET                    84548 non-null  object
17  YEAR BUILT                            84548 non-null  int64
18  TAX CLASS AT TIME OF SALE              84548 non-null  int64
19  BUILDING CLASS AT TIME OF SALE         84548 non-null  object
20  SALE PRICE                             84548 non-null  object
21  SALE DATE                             84548 non-null  object
dtypes: int64(10), object(12)
memory usage: 14.2+ MB
```

```
In [5]: 1 #Check for duplicated data
```

```
In [6]: 1 df.duplicated()
```

```
Out[6]: 0      False
        1      False
        2      False
        3      False
        4      False
        ...
        84543 False
        84544 False
        84545 False
        84546 False
        84547 False
        Length: 84548, dtype: bool
```

```
In [7]: 1 print(df.duplicated().value_counts())
```

```
False      84548
dtype: int64
```

```
In [8]: 1 df=df.drop_duplicates()
```

```
In [9]: 1 print(df.duplicated().value_counts())
```

```
False      84548
dtype: int64
```

```
In [10]: 1 #Check for missing data
```


```
In [11]: 1 df.isnull().sum()
```

```
Out[11]: Unnamed: 0      0
        BOROUGH      0
        NEIGHBORHOOD  0
        BUILDING CLASS CATEGORY  0
        TAX CLASS AT PRESENT  0
        BLOCK      0
        LOT      0
        EASE-MENT  0
        BUILDING CLASS AT PRESENT  0
        ADDRESS      0
        APARTMENT NUMBER  0
        ZIP CODE      0
        RESIDENTIAL UNITS  0
        COMMERCIAL UNITS  0
        TOTAL UNITS      0
        LAND SQUARE FEET  0
        GROSS SQUARE FEET  0
        YEAR BUILT      0
        TAX CLASS AT TIME OF SALE  0
        BUILDING CLASS AT TIME OF SALE  0
        SALE PRICE      0
        SALE DATE      0
        dtype: int64
```

```
In [12]: 1 import missingno as msno
```

```
In [13]: 1 import matplotlib.pyplot as plt
```

```
In [14]: 1 msno.matrix(df, fontsize=12, labels=True)
2 plt.title('Missing data')
3 plt.show()
```



```
In [15]: 1 #Convert variables to a proper format and drop hidden null values among quant
```

```
In [16]: 1 df['LAND SQUARE FEET'] = pd.to_numeric(df['LAND SQUARE FEET'],
2 errors='coerce')
```

```
In [17]: 1 df = df.dropna(subset=['LAND SQUARE FEET'])
```

```
In [18]: 1 df['GROSS SQUARE FEET'] = pd.to_numeric(df['GROSS SQUARE FEET'],
2 errors='coerce')
```

```
In [19]: 1 df = df.dropna(subset=['GROSS SQUARE FEET'])
```

```
In [20]: 1 df['SALE PRICE'] = pd.to_numeric(df['SALE PRICE'], errors='coerce')
```

```
In [21]: 1 df = df.dropna(subset=['SALE PRICE'])
```


```
In [ ]: 1
```

In [22]: 1 df.head()

Out[22]:

Unnamed: 0	BOROUGH	NEIGHBORHOOD	BUILDING CLASS CATEGORY	TAX CLASS AT PRESENT	BLOCK	LOT	EASE- MENT	BUILDING CLASS CATEGORY
0	4	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	392	6	
3	7	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	402	21	
4	8	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	404	55	
6	10	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	406	32	
9	13	1	ALPHABET CITY	08 RENTALS - ELEVATOR APARTMENTS	2	387	153	

5 rows × 22 columns



In [ ]: 1

In [ ]: 1

In [23]: 1 #####

In [24]: 1 #Fix variables so that they're in integer format

In [25]: 1 df.BOROUGH.unique()

Out[25]: array([1, 2, 3, 4, 5], dtype=int64)

```
In [26]: 1 df['RESIDENTIAL UNITS'].unique()
```

```
Out[26]: array([ 5, 10, 6, 8, 24, 3, 4, 0, 1, 2, 22,
 9, 15, 30, 35, 11, 28, 7, 18, 12, 16, 20,
21, 19, 48, 529, 476, 317, 14, 42, 107, 31, 36,
34, 33, 74, 29, 23, 25, 286, 256, 771, 257, 38,
311, 41, 51, 76, 45, 72, 56, 68, 37, 50, 26,
17, 61, 60, 894, 67, 184, 78, 181, 13, 102, 121,
52, 27, 70, 369, 77, 40, 233, 91, 32, 109, 387,
153, 193, 62, 146, 94, 46, 44, 84, 75, 65, 95,
49, 63, 64, 100, 54, 43, 716, 680, 47, 179, 71,
39, 53, 55, 283, 66, 114, 59, 89, 73, 79, 83,
498, 81, 162, 127, 222, 99, 130, 90, 58, 159, 462,
142, 948, 129, 889, 271, 150, 120, 57, 117, 106, 85,
103, 118, 164, 139, 165, 122, 113, 134, 278, 135, 324,
180, 96, 144, 152, 88, 136, 291, 1844, 198, 148, 446,
335, 143, 128], dtype=int64)
```

```
In [27]: 1 df['RESIDENTIAL UNITS_numeric']=df['RESIDENTIAL UNITS']
```

```
In [28]: 1 dict_res={'RESIDENTIAL UNITS_numeric':{'5.': '5', '10.': '10', '6.': '6',
2          '8.': '8',
3          '1.': '1', '3.': '3',
4          '4.': '4', '0.': '0', '2.': '2',
5          '9.': '9', '15.': '15', '11.': '11',
6          '7.': '7',
7          '18.': '18', '12.': '12',
8          '16.': '16', '20.': '20',
9          '19.': '19', '14.': '14', '17.': '17',
10         '13.': '13'}}}
```

```
In [29]: 1 df.replace(dict_res, inplace=True)
```

```
In [30]: 1 df['COMMERCIAL UNITS'].unique()
```

```
Out[30]: array([ 0, 1, 2, 3, 13, 5, 4, 19, 10, 14, 8,
 6, 35, 55, 17, 12, 15, 9, 23, 52, 318, 11,
254, 7, 26, 59, 62, 42, 32, 20, 22, 28, 147,
184, 25, 172, 436, 16, 2261, 51, 18, 21, 126],
dtype=int64)
```

```
In [31]: 1 df['COMMERCIAL UNITS_numeric']=df['COMMERCIAL UNITS']
```

```
In [32]: 1 dict_comm={'COMMERCIAL UNITS_numeric':{'0.': '0', '1.': '1', '2.': '2',
2         '3.': '3',
3         '13.': '13', '5.': '5',
4         '4.': '4', '19.': '19',
5         '10.': '10',
6         '14.': '14', '8.': '8', '6.': '6',
7         '17.': '17', '12.': '12',
8         '15.': '15', '9.': '9', '23.': '23',
9         '11.': '11', '7.': '7', '26.': '26',
10        '20.': '20', '22.': '22',
11        '28.': '28', '25.': '25',
12        '16.': '16',
13        '18.': '18', '21.': '21'}}}
```

```
In [33]: 1 df.replace(dict_comm, inplace=True)
```

```
In [34]: 1 df = df.dropna(subset=['GROSS SQUARE FEET'])
```

```
In [35]: 1 df['GROSS SQUARE FEET'] = df['GROSS SQUARE FEET'].replace(' - ', np.nan) # /
```

```
In [36]: 1 mean_value = df['GROSS SQUARE FEET'].mean()
```

```
In [37]: 1 df['LAND SQUARE FEET'].unique()
```

```
Out[37]: array([ 1633.,  2272.,  2369., ..., 11088., 208033., 10796.])
```

```
In [38]: 1 df['LAND SQUARE FEET'] = df['LAND SQUARE FEET'].astype(int)
```

```
In [39]: 1 df['GROSS SQUARE FEET'].unique()
```

```
Out[39]: array([ 6440.,  6794.,  4615., ...,  977.,  2683., 64117.])
```

```
In [40]: 1 df['GROSS SQUARE FEET'] = df['GROSS SQUARE FEET'].astype(int)
```

```
In [41]: 1 #Ordinal Encoding
```

```
In [42]: 1 df['TAX CLASS AT PRESENT'].unique()
```

```
Out[42]: array(['2A', '2B', '2', '4', '1', '2C', '1A', '1B', '3', ' ', '1C'],
              dtype=object)
```

```
In [43]: 1 df['TAX_numeric']=df['TAX CLASS AT PRESENT']
```

```
In [44]: 1 dict_tax={'TAX_numeric':{' ':0, '1':1, '1A':2,
2         '1B':3, '1C':4, '2':5, '2A':6,
3         '2B':7, '2C':8, '3':9, '4':10}}
```

```
In [45]: 1 df.replace(dict_tax, inplace=True)
```

In [ ]:

1

In [ ]:

1

In [ ]:

1

In [ ]:

1

In [46]:

1 *#Check for outliers*

In [47]:

1 **import** seaborn

In [48]:


1 **from** pandas **import** DataFrame

In [49]:


1 **import** scipy.stats **as** stats

In [50]:


1 boxplot1=seaborn.boxplot(x='RESIDENTIAL UNITS', data=df)



```
In [51]: 1 boxplot2=seaborn.boxplot(x='COMMERCIAL UNITS', data=df)
```

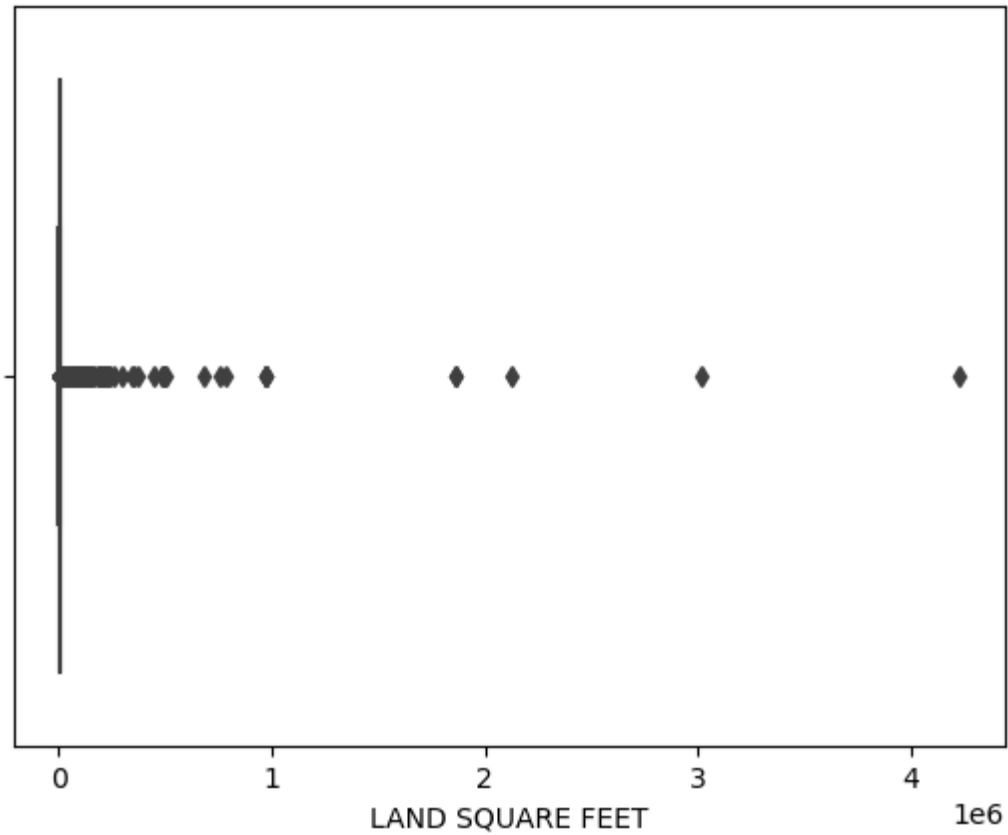


```
In [52]: 1 boxplot3=seaborn.boxplot(x='TOTAL UNITS', data=df)
```







```
In [53]: 1 boxplot4=seaborn.boxplot(x='LAND SQUARE FEET',data=df)
```



```
In [54]: 1 boxplot5=seaborn.boxplot(x='GROSS SQUARE FEET',data=df)
```



```
In [55]: 1 boxplot6=seaborn.boxplot(x='SALE PRICE', data=df)
```




```
In [ ]: 1
```

```
In [56]: 1 #Fix outliers
```

```
In [57]: 1 median=np.median(df['RESIDENTIAL UNITS'])
```

```
In [58]: 1 df['RESIDENTIAL UNITS']=np.where(df['RESIDENTIAL UNITS']>20,  
2      median, df['RESIDENTIAL UNITS'])
```

```
In [59]: 1 boxplot_n1=seaborn.boxplot(x='RESIDENTIAL UNITS',data=df)
```




```
In [60]: 1 median2=np.median(df['COMMERCIAL UNITS'])
```

```
In [61]: 1 df['COMMERCIAL UNITS']=np.where(df['COMMERCIAL UNITS']>30,  
2      median2, df['COMMERCIAL UNITS'])
```




```
In [65]: 1 boxplot_n3=seaborn.boxplot(x='TOTAL UNITS', data=df)
```



```
In [66]: 1 median4=np.median(df['LAND SQUARE FEET'])
```

```
In [67]: 1 df['LAND SQUARE FEET']=np.where(df['LAND SQUARE FEET']>4000,  
2      median4, df['LAND SQUARE FEET'])
```

```
In [68]: 1 boxplot_n4=seaborn.boxplot(x='LAND SQUARE FEET', data=df)
```




```
In [69]: 1 median5=np.median(df['GROSS SQUARE FEET'])
```

```
In [70]: 1 df['GROSS SQUARE FEET']=np.where(df['GROSS SQUARE FEET']>4200,  
2      median5, df['GROSS SQUARE FEET'])
```



```
In [74]: 1 boxplot_n6=seaborn.boxplot(x='SALE PRICE', data=df)
```



```
In [ ]: 1
```

```
In [75]: 1 #Export cleaned dataset
```

```
In [76]: 1 df.to_csv('NYC_clean.csv')
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [77]: 1 ##RQ- What characteristics of the units influence the price of the unit?
```

```
In [78]: 1 import statsmodels.formula.api as smf
```

```
In [79]: 1 from sklearn import linear_model
```



```
In [80]: 1 import statsmodels.api as sm
```

```
In [81]: 1 X=df[['BOROUGH','RESIDENTIAL UNITS','COMMERCIAL UNITS',  
2           'LAND SQUARE FEET','GROSS SQUARE FEET','TAX_numeric']]
```

```
In [82]: 1 y=df['SALE PRICE']
```

```
In [83]: 1 model=smf.ols('y~X', data=df)
```

```
In [84]: 1 res=model.fit()
```

```
In [85]: 1 print(res.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:                0.052
Model:                  OLS    Adj. R-squared:           0.052
Method:                 Least Squares    F-statistic:        441.0
Date:                  Wed, 19 Jul 2023    Prob (F-statistic):    0.00
Time:                  00:11:02    Log-Likelihood:       -6.8728e+05
No. Observations:      48244    AIC:                 1.375e+06
Df Residuals:          48237    BIC:                 1.375e+06
Df Model:               6
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.811e+05	8371.894	21.638	0.000	1.65e+05	1.98e+05
X[0]	8.795e+04	2054.530	42.806	0.000	8.39e+04	9.2e+04
X[1]	7059.6424	1056.201	6.684	0.000	4989.474	9129.810
X[2]	2530.8511	2364.956	1.070	0.285	-2104.494	7166.196
X[3]	-25.0549	2.065	-12.131	0.000	-29.103	-21.007
X[4]	32.9606	2.196	15.011	0.000	28.657	37.264
X[5]	-6450.9245	749.658	-8.605	0.000	-7920.265	-4981.584

```
=====
Omnibus:                3194.669    Durbin-Watson:           1.421
Prob(Omnibus):           0.000    Jarque-Bera (JB):        3871.056
Skew:                    0.691    Prob(JB):                 0.00
Kurtosis:                 3.125    Cond. No.                 1.40e+04
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.4e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [86]: 1 res.summary()
```

Out[86]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.052
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.052
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	441.0
<b>Date:</b>	Wed, 19 Jul 2023	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	00:11:02	<b>Log-Likelihood:</b>	-6.8728e+05
<b>No. Observations:</b>	48244	<b>AIC:</b>	1.375e+06
<b>Df Residuals:</b>	48237	<b>BIC:</b>	1.375e+06
<b>Df Model:</b>	6		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	1.811e+05	8371.894	21.638	0.000	1.65e+05	1.98e+05
<b>X[0]</b>	8.795e+04	2054.530	42.806	0.000	8.39e+04	9.2e+04
<b>X[1]</b>	7059.6424	1056.201	6.684	0.000	4989.474	9129.810
<b>X[2]</b>	2530.8511	2364.956	1.070	0.285	-2104.494	7166.196
<b>X[3]</b>	-25.0549	2.065	-12.131	0.000	-29.103	-21.007
<b>X[4]</b>	32.9606	2.196	15.011	0.000	28.657	37.264
<b>X[5]</b>	-6450.9245	749.658	-8.605	0.000	-7920.265	-4981.584

<b>Omnibus:</b>	3194.669	<b>Durbin-Watson:</b>	1.421
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	3871.056
<b>Skew:</b>	0.691	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	3.125	<b>Cond. No.</b>	1.40e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.4e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [87]: 1 #Calculate MSE
```

```
In [88]: 1 from sklearn.linear_model import LinearRegression
```

```
In [89]: 1 from sklearn.metrics import mean_squared_error
```

```
In [90]: 1 lin_regress=LinearRegression()
```

```
In [91]: 1 lin_regress.fit(X,y)
```

```
Out[91]: LinearRegression()
```

```
In [92]: 1 y_pred=lin_regress.predict(X)
```

```
In [93]: 1 y_pred
```

```
Out[93]: array([278167.82104213, 291005.00862662, 266787.03549798, ...,  
               577637.87564476, 573865.51527501, 584389.70537782])
```

```
In [94]: 1 mse=mean_squared_error(y,y_pred)
```

```
In [95]: 1 X_=df[['BOROUGH','RESIDENTIAL UNITS','COMMERCIAL UNITS',  
               2 'TAX_numeric','LAND SQUARE FEET','GROSS SQUARE FEET']]
```

```
In [96]: 1 print(mse)
```

```
138470011655.5684
```